

Pose Estimation Using Deep Learning

*A mini project report submitted in partial fulfilment of the requirements for
the award of the degree of*

Bachelor of Technology

in

Computer Science & Engineering

Submitted by

Aman K. Shihab

Aneeta Shajan



Federal Institute of Science And Technology (FISAT)®
Angamaly, Ernakulam

Affiliated to

APJ Abdul Kalam Technological University
CET Campus, Thiruvananthapuram

July 2022

Federal Institute of Science And Technology (FISAT)[®]
Mookkannoor(P.O), Angamaly-683577



CERTIFICATE

This is to certify that the report entitled “**Pose Estimation Using Deep Learning**” is a bonafide record of the mini project submitted by **Aman K. Shihab(FIT19CS015)**, in partial fulfilment of the requirements for the award of the degree of Bachelor of Technology (B.Tech) in Computer Science & Engineering during the academic year 2021-22.

Staff in Charge

Project Guide

Dr. Jyothish K John
Head of the Department

ABSTRACT

Single-person human pose estimation facilitates markerless movement analysis in sports, as well as in clinical applications. Still, state-of-the-art models for human pose estimation generally do not meet the requirements of real-life applications. The proliferation of deep learning techniques has resulted in the development of many advanced approaches. However, with the progresses in the field, more complex and inefficient models have also been introduced, which have caused tremendous increases in computational demands. To cope with these complexity and inefficiency challenges, we propose a novel convolutional neural network architecture, called EfficientPose, which exploits recently proposed EfficientNets in order to deliver efficient and scalable single-person pose estimation. EfficientPose is a family of models harnessing an effective multi-scale feature extractor and computationally efficient detection blocks using mobile inverted bottleneck convolutions, while at the same time ensuring that the precision of the pose configurations is still improved. Due to its low complexity and efficiency, EfficientPose enables real-world applications on edge devices by limiting the memory footprint and computational cost. The results from our experiments, using the challenging MPII single-person benchmark, show that the proposed EfficientPose models substantially outperform the widely-used OpenPose model both in terms of accuracy and computational efficiency. In particular, our top-performing model achieves state-of-the-art accuracy on single-person MPII, with low-complexity ConvNets.

Contribution by Author

Author Contribution Goes Here

Student Name

ACKNOWLEDGEMENT

Your Acknowledgement Goes Here

Student 1

Contents

List of Figures	v
List of Tables	vi
1 Introduction	1
1.1 Overview	1
1.2 Problem Statement	1
1.3 Objective	2
2 Related works	3
3 Design	5
3.1 Introduction	5
3.2 Architecture	5
3.3 Modules	6
3.3.1 EfficientNet	6
3.3.2 E-Swish	6
3.4 Accuracy Measures	6
4 Datasets	7
4.1 MPII Human Pose Dataset	7
4.2 Leeds Sports Dataset	7
5 Results	9
6 Outputs	10
7 Conclusion	13
Appendices	15

List of Figures

1.1	Human Pose Estimation Demo	1
2.1	OpenPose	4
3.1	EfficientPose	5
5.1	Convergence of mean error on OpenPose and EfficientPose	9
6.1	Output 1	10
6.2	Output 1	11
6.3	Output 1	12

List of Tables

5.1	Comparison of EfficientPose and OpenPose on MPII validation dataset	9
-----	---	---

Chapter 1

Introduction

1.1 Overview

Single-person human pose estimation (HPE) refers to the computer vision task of localizing human skeletal keypoints of a person from an image or video frames. Single- person HPE has many real-world applications, ranging from outdoor activity recognition and computer animation to clinical assessments of motor repertoire and skill practice among professional athletes. The proliferation of deep convolutional neural networks (ConvNets) has advanced HPE and further widen its application areas. ConvNet-based HPE with its increasingly complex network structures, combined with transfer learning, is a very challenging task. However, the availability of high-performing ImageNet backbones, together with large tailor-made datasets, such as MPII for 2D pose estimation, has facilitated the development of new improved methods to address the challenges. An increasing trend




Figure 1.1: Human Pose Estimation Demo

in computer vision has driven towards more efficient models. Recently, EfficientNet was released as a scalable ConvNet architecture, setting benchmark record on ImageNet with a more computationally efficient architecture.

1.2 Problem Statement

Despite the availability of more performant, efficient layers and architecture's, it has not quite translated into fruitful results within human pose estimation, there is still a lack of architectures that are both accurate and computationally efficient

at the same time. In general, current state-of-the-art architectures are computationally expensive and highly complex, thus making them hard to replicate, cumbersome to optimize, and impractical to embed into real-world applications.

Minimal and efficient model architectures are coveted for their ability to run on edge devices with minimal hardware requirements. This also greatly reduces the response times thereby becoming increasingly relevant for real-time applications. The ability to run real-time is a crucial feature for pose estimation, because often times in most applications it's desirable to obtain the outputs instantly to judge the output.

1.3 Objective

Primary objective is to exploit recent advances in ConvNets and shed some light onto an improved approach called EfficientPose. The main idea is to modify OpenPose, a well-known pose estimation model into a family of scalable ConvNets for high-precision and computationally efficient single-person pose estimation from 2D images. Then we evaluate the EfficientPose model by comparing it against the original OpenPose model on single-person HPE. After that, we compare it against the current state-of-the-art single-person HPE methods on the official MPII challenge, focusing on accuracy as a function of the number of parameters. EfficientPose models aim to elicit high computational efficiency, while bridging the gap in availability of high-precision HPE networks.

Chapter 2

Related works

Ever since the increased adoption of ConvNets for HPE following the success of DeepPose has set the path for accurate HPE. Another breakthrough in HPE was provided by OpenPose. OpenPose comprises a multi-stage architecture performing a series of detection passes. Provided an input image of 368×368 pixels, OpenPose utilizes an ImageNet pretrained VGG-19 backbone to extract basic features. The features are supplied to a DenseNet-inspired detection block arranged as five dense blocks, each containing three 3×3 convolutions with PReLU activations. The detection blocks are stacked in a sequence. First, four passes of part affinity fields map the associations between body keypoints. Subsequently, two detection passes predict keypoint heatmaps to obtain refined keypoint coordinate estimates. In terms of level of detail in the keypoint coordinates, OpenPose is restricted by its output resolution of 46×46 pixels. The OpenPose architecture can be improved by recent advancements in ConvNets, as follows: First, automated network architecture search has found backbones that are more precise and efficient in image classification than VGG and ResNets. Compound model scaling can balance the image resolution, width (number of networkchannels), and depth (number of network layers). This resulted in scalable convolutional neural networks, called EfficientNets, with which the main goal was to provide lightweight models with a sensible trade-off between model complexity and accuracy across various computational budgets. For each model variant EfficientNet, from the most computationally efficient one being EfficientNet-B0 to the most accurate model, EfficientNet-B7. The total no. of FLOPS increases by a factor of 2, given by:

$$(\alpha.\beta^2.\gamma^2)^\phi$$

where

$$\alpha = 1.2, \beta = 1.1, \gamma = 1.15$$

They denote the coefficients for depth, width and resolution respectively.

Second, parallel multi-scale feature extraction has improved the precision levels in HPE, emphasizing both high spatial resolution and low-scale semantics. However, existing multi-scale approaches in HPE are computationally expensive, both due to their large size and high computational requirements. For example, a typical multi-scale HPE approach has often a size of 16 to 58 million parameters and requires 10 to 128 GFLOPS. To cope with this, we propose cross-resolution features, operating on high- and low-resolution input images, to integrate features from multiple abstraction levels with low overhead in network complexity and with high computational efficiency. Existing works on Siamese ConvNets have been promising in utilizing parallel network backbones. Third, mobile inverted bottleneck convolution (MBConv) with built-in squeeze-and- excitation (SE) and Swish activation integrated in EfficientNets has proven more accurate in image classification




Figure 2.1: OpenPose

tasks than regular convolutions, while substantially reducing the computational costs. The efficiency of MBConv modules stem from the depthwise convolutions operating in a channel-wise manner. With this approach, it is possible to reduce the computational cost by a factor proportional to the number of channels. Hence, by replacing the regular 3×3 convolutions with up to 384 input channels in the detection blocks of OpenPose with MBConvs, we can obtain more computationally efficient detection blocks. Further, SE selectively emphasizes discriminative image features, which may reduce the required number of convolutions and detection passes by providing a global perspective on the estimation task at all times. Using MBConv with SE may have the potential to decrease the number of dense blocks in OpenPose. Fourth, transposed convolutions with bilinear kernel scale up the low-resolution feature maps, thus enabling a higher level of detail in the output confidence maps. The main advantage of this is that we can use ConvNets that are small and computationally efficient enough to run on edge devices with little memory and low processing power, which is impossible with OpenPose. We can also alter the parameters of EfficientNet to obtain different variants with accuracies and efficiencies that are different.

Chapter 3

Design

3.1 Introduction

Here, the model architecture we use, namely EfficientPose exploits the recent advancements in ConvNets and additionally concatenates high-level and low-level features. The net result is a much more efficient model with better accuracy and needing fewer computational resources.

3.2 Architecture




Figure 3.1: EfficientPose

The model takes in two inputs, one is the high-level features and second is the low-level features. The low level features are obtained by downsampling an image to half of it's height and width using an average pooling layer. The feature extractors here are initial layers of EfficientNet. For high level extractor, $\phi \in [0, 7]$ and for the low level extractor $\phi \in [0, 3]$.

These extracted features are then concatenated together to obtain cross-resolution features. This helps us to emphasize the important local factors in the image of interest.

The input to the next phase of the model is the cross-resolution features obtained in the previous step. Here the required keypoints are localized through

an iterative process, where each detection pass performs supervised prediction of output maps. Each detection pass comprises a detection block and a single 1×1 convolution for output prediction. The detection blocks across all detection passes exhibit the same basic architecture, comprising Mobile DenseNets. Data from here is forwarded to the successive layers through skip connections. Here, we also avoid downsampling of the output to preserve the resolution. The original mobile convnets are modified by adding an E-swish activation function with β value 1.25.

The overall detection is performed in two rounds. Initially, the overall pose of the person is anticipated through a single pass of skeleton estimation. This helps especially when there are multiple people present in the frame. After the skeleton estimation is done two detection passes are performed to estimate heatmaps for points of interest.

Another improvement on top of OpenPose is that, EfficientPose projects lower-resolution image onto a higher resolution space using transposed convolution to allow an increased level of detail, whereas in OpenPose the heatmaps are constrained to the lower space.

3.3 Modules

3.3.1 EfficientNet

In mathematics, Stirling's approximation (or Stirling's formula) is an approximation for large factorials. It is named after James Stirling.

3.3.2 E-Swish

3.4 Accuracy Measures

Chapter 4

Datasets

4.1 MPII Human Pose Dataset

MPII Human Pose dataset is a state of the art benchmark for evaluation of articulated human pose estimation. The dataset includes around 25K images containing over 40K people with annotated body joints. The images were systematically collected using an established taxonomy of every day human activities. Overall the dataset covers 410 human activities and each image is provided with an activity label. Each image was extracted from a YouTube video and provided with preceding and following un-annotated frames. In addition, for the test set we obtained richer annotations including body part occlusions and 3D torso and head orientations. The model was trained primarily on this dataset.

4.2 Leeds Sports Dataset

This dataset contains 2000 pose annotated images of mostly sports people gathered from Flickr using the tags shown above. The images have been scaled such that the most prominent person is roughly 150 pixels in length. Each image has been annotated with 14 joint locations. Left and right joints are consistently labelled from a person-centric viewpoint. Attributions and Flickr URLs for the original images can be found in the JPEG comment field of each image file. The ordering of the joints are as follows:

1. Right Ankle
2. Right Knew
3. Right Hip
4. Left Hip
5. Left Knee
6. Left Ankle
7. Right Wrist
8. Right Elbow
9. Right Shoulder
10. Left Shoulder

11. Left Elbow
12. Left Wrist
13. Neck
14. Head Top

Chapter 5

Results




Figure 5.1: Convergence of mean error on OpenPose and EfficientPose

Some text

Table 5.1: Comparison of EfficientPose and OpenPose on MPII validation dataset

Model	Parameters	Parameter Reduction	FLOPs	FLOP Reduction	$PCK_h@50$	$PCK_h@50$
1	China	1,347,350,000	19.24%			
2	India	1,210,193,422	17.28%			
3	United States	313,269,000	4.47%			

Chapter 6

Outputs



Figure 6.1: Output 1




Figure 6.2: Output 1




Figure 6.3: Output 1

Chapter 7

Conclusion

An intrusion detection system (IDS) [1] is a device or software application that monitors network and/or system activities for malicious activities or policy violations and produces reports to a Management Station.

Donald Ervin Knuth [2] is a computer scientist and Professor Emeritus at Stanford University. He is the author of the seminal multi-volume work The Art of Computer Programming. Knuth has been called the "father" of the analysis of algorithms

Bibliography

- [1] K. Scarfone and P. Mell, “Guide to intrusion detection and prevention systems (idps),” *NIST Special Publication*, vol. 800, no. 2007, p. 94, 2007.
- [2] Wikipedia, “Donald knuth.” http://en.wikipedia.org/wiki/Donald_Knuth.

Appendices

Test